

Data Mining: Learning from Large Data Sets - Spring Semester 2014

member1@student.ethz.ch
member2@student.ethz.ch
member3@student.ethz.ch

March 2, 2014

Approximate near-duplicate search using Locality Sensitive Hashing

Briefly describe the steps used to produce the solution. Feel free to add plots or screenshots if you think it's necessary. The report should contain a maximum of 2 pages.